

## CS 475: Concurrent & Distributed Systems

Prof. Sanjeev Setia  
Computer Science Dept  
George Mason University

### About this Class

- ❑ Focus: designing and writing moderate-sized concurrent and distributed applications
  - Fundamental concepts
  - Multi-threaded and distributed programs
- ❑ See syllabus for course learning outcomes
- ❑ Prerequisites:
  - CS 367 (Computer Systems & Programming)
  - High level of competence in C and Java

## What you will learn

*"I hear and I forget, I see and I remember, I do and I understand" - Chinese proverb*

- ❑ Fundamental concepts in the development of concurrent & distributed software
- ❑ Developing Concurrent Programs
  - Threads, semaphores, condition variables, monitors
- ❑ Middleware technology for distributed applications
  - Network programming using TCP/IP Sockets
  - RPC/RMI
    - Web Services

CS 475

3

## Logistics

- ❑ Grade: 50% projects and homework, 50% exams
  - 15% quizzes, 15% midterm, 20% final (cumulative)
  - Tentative date of midterm - Oct 23
- ❑ Four programming assignments and (at least) one paper and pencil
  - First three assignments roughly same weight, fourth higher weight
  - Can be done in groups of two
  - First three assignments use C, fourth Java
  - Assignments will be graded on VS&E Linux server (zeus)
    - If you do your development elsewhere, your responsibility to make sure it runs correctly on zeus
- ❑ Homework problems
  - To be done individually

CS 475

4

## Logistics cont'd

- ❑ Online Assignment submission
  - Blackboard (mymason.gmu.edu)
  - Grades posted on Blackboard
- ❑ Lateness Policy
  - Four "slip" days collectively for four programming assignments
  - Can use at most two slip days for an assignment
  - Late submissions not accepted for "paper and pencil" homework assignments
- ❑ Honor Code
- ❑ Classroom Policy: Use of laptops/PDAs not permitted

CS 475

5

## Logistics cont'd

- ❑ Office Hrs
  - Thursday, 1-3 pm, Room 4300, Engineering Bldg
- ❑ Email: [setia@gmu.edu](mailto:setia@gmu.edu)
- ❑ GTA: Nusha Mehmanesh
  - Email [nmehmane@gmu.edu](mailto:nmehmane@gmu.edu)
  - Office hrs: Monday 1-3 PM, Tuesday 2-4 PM, Wednesday 10 AM - 12 PM
  - Office: TBA
- ❑ Piazza - discussion forum
- ❑ Class materials will be posted on Blackboard page for course

CS 475

6

## Readings

### ❑ Recommended books

- Computer Systems & Programming (Bryant & O'Halloran) - used in CS 367
- Operating Systems: Three Easy Pieces (Arpaci-Dusseau and Arpaci-Dusseau) - online text
  - <http://pages.cs.wisc.edu/~remzi/OSTEP/>
- Distributed Systems: Concepts & Design (Coulouris et al)

### ❑ Read class slides & notes

CS 475

7

## Programming Assignments

### ❑ Assignment 1: Shell Lab - C

- Topic: Creating and managing concurrent processes

### ❑ Assignment 2: Multithreaded programming Lab - C

- Topic: Concurrent programming, synchronization

### ❑ Assignment 3: Network programming lab - C

- Topic: network programming, multi-threaded programming, synchronization

### ❑ Assignment 4: Calendar Lab - Java

- Topic: RMI, distributed application development

CS 475

8

## Schedule (tentative)

- ❑ Concurrent Programming
- ❑ Process Synchronization
- ❑ Parallel processing on Multicores (introduction)
- ❑ Introduction to Networking
- ❑ Sockets; Application-level network protocols
- ❑ Introduction to distributed systems
- ❑ RPC/RMI
- ❑ Web Services (introduction)
- ❑ And if we have time..... which is unlikely
  - Peer-to-peer computing (introduction)
  - Parallel Programming on multi-computers (introduction)
  - Map Reduce...

CS 475

9

## Hardware Architectures


- ❑ Uniprocessors
- ❑ Shared-memory multiprocessors
- ❑ Distributed-memory multicomputers
- ❑ Distributed systems

CS 475

10

## Processes

- ❑ Def: A *process* is an instance of a running program.
  - One of the most profound ideas in computer science.
  - Not the same as "program" or "processor"
- ❑ Process provides each program with two key abstractions:
  - Logical control flow
    - Each program seems to have exclusive use of the CPU.
  - Private address space
    - Each program seems to have exclusive use of main memory.
- ❑ How are these illusions maintained?
  - Process executions interleaved (multitasking)
  - Address spaces managed by virtual memory system




CS 475

11

## Concurrent Processes

- ❑ Two processes *run concurrently (are concurrent)* if their flows overlap in time.
- ❑ Otherwise, they are *sequential*.
- ❑ Examples:
  - Concurrent: A & B, A & C
  - Sequential: B & C



CS 475

12

## Cooperating Concurrent Processes

- ❑ Concurrent processes part of the same application
- ❑ Processes "cooperate" on task
- ❑ Motivation
  - Support inherent concurrency in application
    - Window systems, web servers
  - Improved performance - can make use of multiple processors

CS 475

13

## Concurrent Programming

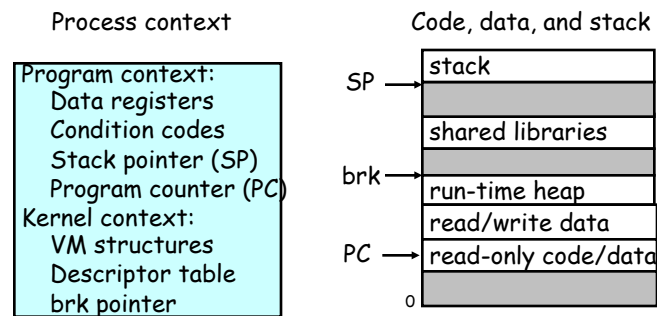
- ❑ Process = Address space + one thread of control
- ❑ Concurrent program = **multiple threads of control**
  - Multiple single-threaded processes
  - Multi-threaded process

CS 475

14

## Traditional View of a Process

- Process = process context + code, data, and stack

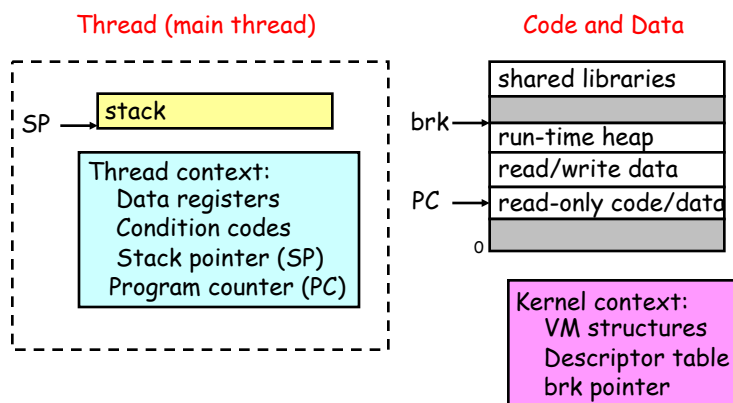


CS 475

15

## Alternate View of a Process

- Process = thread + code, data, and kernel context



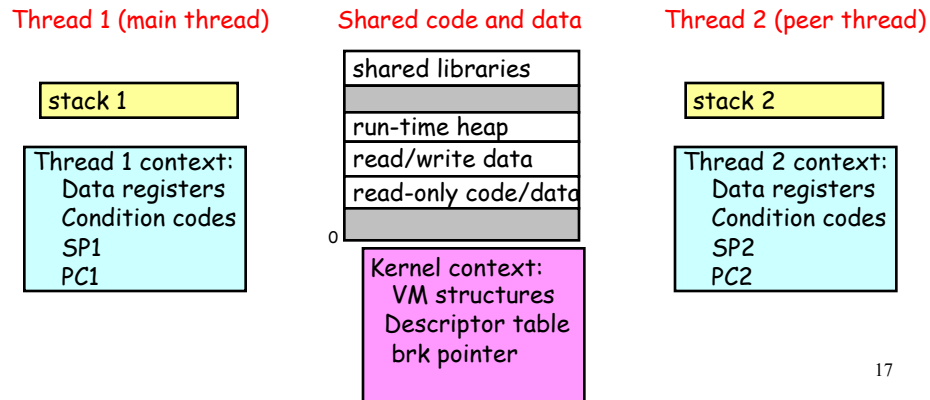
CS 475

16



## A Process With Multiple Threads

- ❑ Multiple threads can be associated with a process
  - Each thread has its own logical control flow (sequence of PC values)
  - Each thread shares the same code, data, and kernel context
  - Each thread has its own thread id (TID)



17

## Threads: Motivation

- ❑ Traditional processes created and managed by the OS kernel
- ❑ Process creation expensive - fork system call in UNIX
- ❑ Context switching expensive
- ❑ Cooperating processes - no need for memory protection (separate address spaces)

CS 475

18

## Threads

- ❑ Execute in same address space
  - separate execution stack, share access to code and (global) data
- ❑ Smaller creation and context-switch time
- ❑ Can exploit fine-grain concurrency

CS 475

19

## Challenges in multi-threaded/concurrent programming

- ❑ Synchronizing multiple processes/threads
  - Locks
  - Semaphores
  - Monitors
  - Deadlocks
  - Livelocks
- ❑ Testing/debugging concurrent applications is a lot harder!

CS 475

20

## Application classes

- ❑ Multi-threaded Programs
  - Processes/Threads on same computer
  - Window systems, Operating systems
- ❑ Distributed computing
  - Processes/Threads on separate computers
  - File servers, Web servers
- ❑ Parallel computing
  - On same (multiprocessor) or different computers
  - Goal: solve a problem faster or solve a bigger problem in the same time

CS 475

21

## Distributed systems

- ❑ "Workgroups"/Intranets
- ❑ ATM (bank) machines
- ❑ World wide web
- ❑ Multimedia conferencing
- ❑ Ubiquitous network-connected devices
  - Internet of Things

CS 475

22

## Distributed applications

- ❑ Applications that consist of a set of processes that are distributed across a network of machines and work together as an ensemble to solve a common problem
- ❑ In the past, mostly “client-server”
  - Resource management centralized at the server
- ❑ Peer-to-peer applications represent “truly” distributed applications

CS 475

23

## Goals/Benefits

- ❑ Resource sharing
- ❑ Scalability
- ❑ Fault tolerance and availability
- ❑ Performance
  - Parallel computing can be considered a subset of distributed computing

CS 475

24

## Challenges (Differences from Local Computing)

- ❑ Heterogeneity
- ❑ Latency
  - Interactions between distributed processes have a higher latency
- ❑ Memory Access
  - Remote memory access is not the same as local memory access
    - Local pointers are meaningless outside address space of process

CS 475

25

## Challenges cont'd

- ❑ Synchronization
  - Concurrent interactions the norm
- ❑ Partial failure
  - Applications need to adapt gracefully in the face of partial failure
  - Leslie Lamport (a famous computer scientist) once defined a distributed system as "One on which I cannot get any work done because some machine I have never heard of has crashed"

CS 475

26


## Communication Patterns

- ❑ Client-server
- ❑ Group-oriented
  - Applications that require reliability
- ❑ Function-shipping
  - Java applets

CS 475

27


## Clients invoke individual servers



CS 475

28


### A service provided by multiple servers



CS 475

29


### Web proxy server



CS 475

30

## A distributed application based on peer processes




CS 475

31

## Web applets

a) client request results in the downloading of applet code



b) client interacts with the applet




CS 475

32



## Thin clients and compute servers

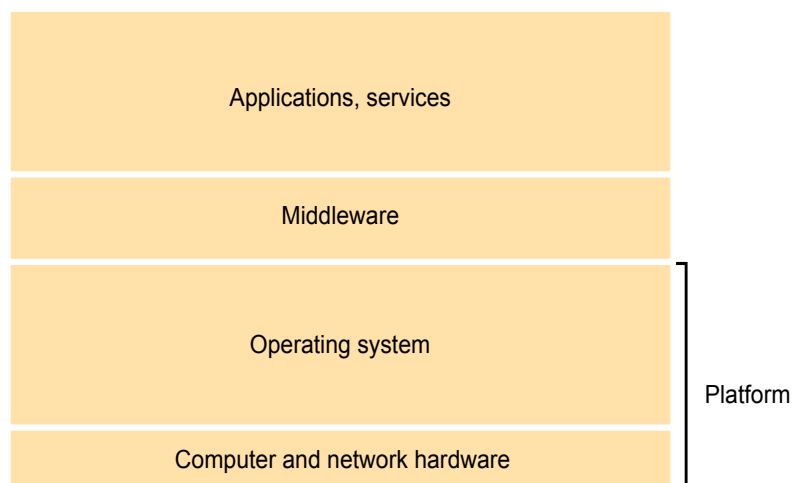


Cloud Computing

CS 475

33

## Software and hardware service layers in distributed systems



CS 475

34

## Road Map

- ❑ Next class: Processes and Threads
- ❑ Next week: Processes and Signals
  - Continuation of material introduced in CS 367
- ❑ Week 3-5: Concurrent programming
- ❑ Assignment 1: Shell Lab (officially assigned next week but may be posted on Blackboard this week)